

Adhithya Rokhith Bhaskar

Boston, MA | arokhith@gmail.com | (929) 563-9974

May 8, 2026

Hiring Team, The Nielsen Company

Full-funnel digital campaign management at scale breaks at one specific seam: the gap between the platform-native dashboards each channel team owns and the multi-touch attribution view that explains how those channels actually compound on a converted account. Most teams have respectable channel-level numbers, a GA4 instance that broadly works, and a quarterly executive review that shows blended performance, and none of those reconcile against the simple question "which combination of touches actually drove the closed-won." The campaign managers who survive at media-insights scale are the ones who have rebuilt the loop with one UTM spec, one source-of-truth, a real attribution model, and an executive narrative the leadership team can act on by Friday.

I have done that rebuild. At Plug Power I owned full-funnel digital campaigns across awareness, engagement, and conversion: Google Ads, Meta, LinkedIn, and Adobe-driven programmatic with -25% CPA and +20% conversion. The multi-touch attribution rebuild ran across HubSpot, Salesforce, Adobe Analytics, GA4, and GTM. UTM spec written from scratch. GTM containers rebuilt. Lead routing validated through Salesforce campaign attribution before any cohort number got presented. Pixel and tag hygiene was the foundation, not an afterthought.

On the executive-narrative side, Pixis is the analog. Eight-client B2B SaaS portfolio with per-client monthly media budgets ranging \$15K to \$90K, weekly client CFO QBRs presenting ROI and iROAS to growth leads and finance heads. Replaced two underperforming channels mid-quarter when KPIs slipped. The "translate metrics into actionable executive narratives" line in the JD is the literal description of those QBRs.

On the A/B-testing framework, SwipeHire shipped solo on Python, OpenAI API, and React with 8 A/B tests on activation flows tied back to Amplitude cohort analysis and a +25% activation lift. Pixis added 3x ROAS validation via holdout incrementality. Plug Power layered keyword, audience, and creative tests across the paid stack. Three different scales of testing rigor, same methodology.

On AI-native execution, four Anthropic Claude certifications come from running prompt-based iteration as my actual workflow. A Claude draft-and-critique pipeline I productized delivers 8x velocity on copy variants.

Remote works cleanly. Resume attached.

Portfolio: adhi-growth-ai.vercel.app/ref/nielsen-dcm

Adhithya Rokhith Bhaskar