

Adhithya Rokhith Bhaskar

Boston, MA | arokhith@gmail.com | (929) 563-9974

April 27, 2026

Hiring Team, OpenAI

Rapid response social listening breaks at one specific place: the gap between a recurring narrative and the moment it crosses from observation into action. Most "we missed it" stories are not detection failures. They are clustering and velocity failures. Three sub-narratives that look different in isolation are often one story moving fast, and the team that catches it is the one that has built the clustering taxonomy and the velocity threshold before there is a fire.

I have done that work at smaller scale, in the same forums OpenAI's analyst would monitor. SwipeHire is an AI product I built solo on Python, OpenAI API, and React, and I ran organic community presence inside Reddit (r/cscareerquestions, r/h1b), Twitter threads, and weekly product updates. Living inside those communities for a year is how I learned that the early signal of a brewing AI-product narrative is almost always a small cluster of high-engagement Reddit comments before it shows up anywhere else. Discord and Hacker News follow with a 24 to 72 hour lag.

On signal-from-noise, the cleanest example came at Pixis. Two enterprise clients on an 8-client portfolio were ready to leave. The aggregate ROAS dashboards looked fine. I pulled session-level data and found creative fatigue masked by blended metrics: the top 5% of sessions were carrying the average. Refresh the creative, re-key the channel mix, weekly health checks. Both clients retained. The lesson generalizes: aggregate metrics hide the signal that matters.

On dashboard-build muscle, I designed and presented a \$627M margin-recovery executive Power BI dashboard at Sensata tracking 1,200 opportunities across three regions, four product families, and 20 product lines. Bleeders versus leakers framing, regional scorecards, win-rate analytics. Senior leadership read it weekly. The same approach maps to narrative-cluster dashboards: classify, score velocity, surface the ones crossing escalation thresholds.

On qualitative-to-structured translation, I ran 30+ customer interviews at Sensata with OEMs and distributors and turned the unstructured conversation into product themes and CRO experiments that raised CSAT 22%. That same translation muscle is what this role calls "translate online conversations into concise insights for internal stakeholders."

SF hybrid is fine. Happy to walk through any of the above in more detail.

Portfolio: adhi-growth-ai.vercel.app/ref/openai-social

Adhithya Rokhith Bhaskar