

# Adhithya Rokhith Bhaskar

Boston, MA | arokhith@gmail.com | (929) 563-9974

May 18, 2026

Hiring Team, Amazon News

Organic search is being rewritten by LLMs in real time, and AboutAmazon.com is one of the rare newsrooms with enough domain authority, editorial cadence, and structural variety to lead through that rewrite rather than chase it. The role you've described - one seat owning both organic search and AI visibility, with technical-SEO and schema chops sitting next to first-person LLM mechanics and a content-first editorial partnership - is the exact intersection I've been operating at for the last four years, so I'm writing.

On SEM + SEO + content together, Plug Power is the closest analog. I rebuilt the digital marketing function for a B2B hydrogen-energy audience from zero: paid + organic on the same content-cluster map, A/B tests on keywords, landing-page copy, meta + headline structure across paid and organic, and the technical SEO floor underneath - canonical tags, schema markup, internal linking, Core Web Vitals, page-speed work, all on WordPress. CPA fell 25%, conversion rose 20%, and the SEMrush + Google Search Console competitor benchmarking became the cluster map the editorial team published against. I followed that with an Enterprise SEO Gap Analysis as a deep portfolio project: technical audit (crawl, Core Web Vitals, schema), content-cluster mapping, competitor keyword benchmarking, and a prioritized execution roadmap built for in-house ownership.

On structured data and feed-level optimization, GenY Medium is the proof. I rebuilt a D2C e-commerce brand's Google Merchant Center feeds end-to-end - fixed disapprovals, restructured Shopping campaigns with feed-driven titles, and shipped schema markup and structured product data to Google. That feed work, paired with founder-facing weekly performance dashboards, drove a 16x ROAS rebuild. Structured data is not a side discipline for me; it's where I learned that schema is what makes a corpus legible to anything that reads code-first - which is exactly the muscle GEO asks for now.

On LLM content surfacing and AI discoverability, I have hands-on builder credentials: SwipeHire is a production AI matching product I shipped solo on Python + OpenAI API + React with day-one GA4 + GTM event taxonomy, 8 A/B tests with sample-size and significance gates defined up front, and weekly cohort retention review. NSE Paper-Trading Bot is a second AI product I shipped solo on Python + Claude API doing agentic reasoning over news, catalysts, and macro signals. I hold four Anthropic Claude certifications and use Cursor + Claude as a daily workflow. The pattern that builds with LLMs teaches you what LLMs cite: structured context, schema-annotated content, named entities, source-linked claims, and FAQ blocks. The same content that an LLM finds easy to cite is the content that wins citation in Google's AI Overviews. SEO and GEO are not two strategies - they're one discipline at different surfaces.

On translating findings into narratives for leadership, the Sensata \$627M margin-recovery executive Power BI dashboard is the most recent example - 1,200 opportunities across 3 regions and 20 product lines, with a KPI framework, regional scorecards, and top-opportunity focus lists that surfaced \$433M in margin gaps and became a weekly leadership artifact rather than a deck. On cross-functional editorial, comms, and engineering work, Sensata's multi-region GTM across NA and EU is the most direct: positioning, content, and sales enablement sequenced across three regions, generating +35% qualified leads and \$2M+ in first-year pipeline.

NYC works. Boston relocates. The portfolio with the full proof-point trail and a draft of how I'd structure an AboutAmazon.com GEO baseline lives at [adhithyabhaskar.vercel.app/ref/amazon-news](https://adhithyabhaskar.vercel.app/ref/amazon-news). Looking forward to the conversation.

**Adhithya Rokhith Bhaskar**