

PITCH DECK | SEO/GEO MANAGER | AMAZON NEWS

SEO and GEO are one discipline.

Citations - in SERPs, in AI Overviews, in chat answers. Same content spine, different surfaces.

Plug Power - SEM + SEO + content + Mar-Tech rebuild on WordPress (-25% CPA, +20% conversion)


GenY - Google Merchant Center technical rebuild + structured product data (16x ROAS)

SwipeHire - shipped solo on Python + OpenAI API + React (8 A/B tests, +25% activation)

4 Anthropic Claude Certifications + Cursor daily

Enterprise SEO Gap Analysis - crawl, schema, content clusters, competitor benchmarks

NSE Bot - Python + Claude AI + Perplexity + ChatGPT + Anthropic reasoning



one content spine, six citation surfaces

Adhithya Rokhith Bhaskar

Boston, MA | NYC | arokhith@gmail.com | (929) 563-9974 | adhithyabhaskar.vercel.app/ref/amazon-news

Each criterion on the left has a working proof on the right.

Six lines from the posting. Six things already shipped. Substance, not slogan.

What the JD asks for

"Develop comprehensive organic SEO and GEO strategy with growth initiatives."

"Monitor and optimize content visibility in AI chatbots like ChatGPT and Perplexity."

"Conduct technical SEO audits addressing site architecture and indexation."

"Partner with engineering teams on URL structure and schema optimization."

"Build SEO/GEO reporting processes for leadership and stakeholders."

"Work with editorial teams to maximize organic engagement."

What I've actually shipped

Plug Power: SEM + SEO + content together on one WordPress cluster map. Paid and organic A/B tested on the same keywords + landing copy. -25% CPA, +20% conversion. SEMrush + Google Search Console as the competitor-benchmark spine for the editorial team.

Two production AI/LLM products shipped solo + 4 Anthropic Claude certifications. SwipeHire (Python + OpenAI API + React, 50+ users, +25% activation). NSE Paper-Trading Bot (Python + Claude API agentic reasoning over news + catalysts). Cursor + Claude daily workflow.

Enterprise SEO Gap Analysis portfolio project. Crawl, Core Web Vitals, schema audit, content-cluster mapping, competitor keyword benchmarking. Delivered as a prioritized execution roadmap for in-house ownership. Same audit shape AboutAmazon.com would benefit from on day one.

GenY Google Merchant Center technical rebuild. Fixed disapprovals end-to-end. Shipped schema markup + structured product data to Google. 16x ROAS rebuild. Plus Plug Power canonicals, internal linking, page-speed, and Core Web Vitals work on WordPress.

Sensata \$627M Power BI executive dashboard. 1,200 opportunities across 3 regions and 20 product lines, with % to target + win rate + avg \$/win + pipeline-stage KPIs. Became a weekly leadership artifact that surfaced \$433M in margin gaps.

Sensata multi-region GTM (NA + EU + 3 regions). Sequenced positioning + content + sales enablement across engineering, product, and sales. +35% qualified leads, \$2M+ first-year pipeline, time-to-market down 20%.

For AboutAmazon.com, the citation question runs across five query shapes.

A first-week audit looks like this. Score each LLM and AI Overview for whether AboutAmazon.com is the cited source on the queries the newsroom should already own. The pattern of which buckets are weak tells you where the editorial + schema + canonical work should land first.

WHY THIS SHAPE MATTERS

The patterns of weakness tell you which discipline to prioritize. If LLMs miss policy + community + sustainability buckets, the fix is structured-data hygiene (FAQ + Q&A + NewsArticle schema, source-linked claims, canonical clarity) - not more content volume.

That same schema layer is what makes Google AI Overviews and ChatGPT cite the same passages. One spine, six surfaces.

Citation share by query shape - illustrative first-week scoring template					
QUERY BUCKET	ChatGPT	Perplexity	Claude	AI Overview	Top action
Corporate news + product launches "what did Amazon announce at re:Invent 2026"	H	H	H	H	Defend - canonical hygiene
Policy + governance positions "Amazon's position on warehouse worker safety"	M	M	M	L	FAQ + Q&A schema sprint
Leadership + executive narrative "who runs Amazon's AI strategy"	M	H	M	H	Person schema + internal lir
Sustainability + climate commitments "what is Amazon doing about carbon emissions"	L	M	L	M	Source-linked claims overha
Community + philanthropy stories "Amazon disaster-response donations 2025"	L	L	L	M	NewsArticle schema + dates

LEGEND: H = high citation share M = mixed / weak L = AboutAmazon not cited

illustrative pattern - actual scores from week-one audi

Editorial sits on top. Schema, canonicals, and entity hygiene are the spine.

The JD asks for technical SEO audits, schema + URL structure partnership with engineering, and reporting processes for leadership - in one seat. Treated as four layers of a single stack, not four disciplines, the work compounds. Editorial sits where it belongs - on top of a foundation that makes its work easier to find and easier to cite.

LAYER 04 - MEASUREMENT

GSC + SEMrush + LLM-citation tracking + leadership readouts in Power BI / Looker

Impressions, clicks, AI Overview share, LLM-citation share by query bucket - one weekly artifact for the newsroom + leadership

LAYER 03 - EDITORIAL PARTNERSHIP

Cluster mapping + internal linking + named-entity hygiene + FAQ block patterns

Editorial decides the story. Spine makes the story citable. Newsroom playbook: pre-publish checklist + post-publish refresh cycle.

LAYER 02 - STRUCTURED DATA SPINE

Schema sprint: Article + NewsArticle + FAQ + Organization + Person schemas on the top 50 templates

Source-linked claims, canonical tags, internal-link discipline. The layer that makes the corpus legible to anything that reads code-first.

LAYER 01 - TECHNICAL FOUNDATION

Crawl + index + canonicals + Core Web Vitals + URL structure partnership with engineering

Site architecture audit, redirect maps, sitemap hygiene, page-speed budget - the floor everything else stands on.

WHAT EACH LAYER EARNS

Layer 1: SERP credibility. Without crawl + index + canonicals, nothing else compounds.

Layer 2: Citation legibility. Schema + source-linked claims are why LLMs quote *AboutAmazon.com* instead of someone else.

Layer 3: Editorial leverage. The newsroom keeps full ownership of voice; the spine just makes the work easier to find.

Layer 4: Decisions. The dashboard the leadership team opens on Monday - not the deck that sits in a drive.

Three workstreams I'd run in the first thirty days.

Concrete and scoped. Each ships inside a month and produces a measurable artifact - not a deliverables list.

WORKSTREAM 1 - BASELINE

AboutAmazon.com GEO + SEO citation baseline

Score 50 priority queries across ChatGPT, Perplexity, Claude, and Google AI Overviews for whether AboutAmazon.com is the cited source. Pair with a traditional technical-SEO audit on the top 200 templates (crawl, index, schema, Core Web Vitals, canonicals).

Why first. Without the baseline, the rest is opinion. Enterprise SEO Gap Analysis is the same audit shape I've already run end-to-end.

WORKSTREAM 2 - SCHEMA

Schema + structured-data sprint on the top 50 editorial templates

Article + NewsArticle + FAQ + Organization + Person schemas on the top 50 templates. Source-linked claims overhaul on the policy + sustainability buckets. Canonical-tag tightening on the top 25 stories. Engineering RFC for URL structure where indexation is leaking.

Why first. Schema is the spine that wins both SERP and LLM citation - same fix, two surfaces. GenY structured product data is the closest analog I've shipped.

WORKSTREAM 3 - READOUT

Weekly leadership readout cadence + newsroom enablement

Stand up a one-page weekly dashboard: impressions, clicks, AI Overview share, LLM-citation share by query bucket. Plus a one-page weekly note for the newsroom on which story shapes are getting cited by LLMs and which aren't - feeding the next week's editorial cluster decisions.

Why first. Same shape as the Sensata \$627M Power BI executive dashboard - the artifact leadership opens Monday morning, not the deck that sits in a drive.

What I'd want from the first fifteen minutes.

What I'd want to understand: which query buckets the team thinks AboutAmazon.com under-cites today; how the editorial + comms + engineering trio is currently sequenced on schema work; and where the GEO program is most blocked today (tooling, data, or org velocity).

What I'd bring: a draft GEO baseline scoped to 50 priority queries across ChatGPT, Perplexity, Claude, and Google AI Overviews, a 30-60-90 plan for the schema spine, and a one-pager on how I'd structure the weekly leadership readout.

NYC works. Boston relocates. Open to hybrid or onsite cadence.

15 min with the Amazon News team

or written async if that's easier on your week

Adhithya Rokhith Bhaskar

arokhith@gmail.com | (929) 563-9974 | Boston, MA

Portfolio: adhithyabhaskar.vercel.app/ref/amazon-news

LinkedIn: linkedin.com/in/adhithyarokhith