

Adhithya Rokhith Bhaskar

Boston, MA | arokhith@gmail.com | (929) 563-9974

May 13, 2026

Hiring Team, Flagright

The JD reads like a manifesto. "Pipeline quality over CPL." "Cut weak spend, double down on what works." "Push back on activity that doesn't drive pipeline." That's the operating philosophy I've run on for the last three years, and the substance match is unusually clean.

On long-cycle B2B SaaS paid, Pixis is the analog. I ran an 8-client B2B SaaS portfolio at \$1.4M aggregate annual media book - per-client monthly budgets between \$15K and \$90K - across Google, Meta, LinkedIn ABM, Criteo retail-media, Amazon Ads, and programmatic. Long-consideration buyers, multi-month sales cycles, weekly client CFO QBRs translating engagement-to-pipeline. On one beauty client, a holdout-incrementality slice exposed two channels getting credit they hadn't earned - reallocated spend on true lift, brand cleared 3x ROAS the next month. That is exactly the "pipeline quality, not CPL" muscle the JD names. On two other clients, I replaced two underperforming channels mid-quarter when blended numbers hid the truth. The "cut weak spend" reflex isn't a slogan; it's a habit.

On closed-loop attribution and enterprise retargeting, Plug Power is the proof. I rebuilt the MarTech function from zero on Salesforce + Adobe Analytics + GA4 + GTM. UTM spec written from scratch, GTM containers rebuilt, lead-routing rules and SLA documentation feeding BDR and AE alerting. Account-level retargeting funnels built in Adobe Audience Manager with behavioral segments, audience layers by title and seniority and industry, suppression lists. CPA fell 25%, conversion rose 20%, and the loop from ad click to Sales-Qualified Lead got traced and proven.

On landing-page experimentation, SwipeHire is the working example. I built an AI matching product solo on Python + OpenAI API + React, instrumented GA4 + GTM + Amplitude day one, and ran 8 A/B tests on activation flows with sample-size and significance gates defined up front. Activation lifted 25% V1 to V2.

On the culture criteria - knocked down and got back up, resourceful, strong opinions, learn fast, no-BS - the two AI products I shipped solo (SwipeHire on OpenAI API + NSE Paper-Trading Bot on Claude API) are the working proof. Plus four Anthropic Claude certifications and Cursor + Claude as daily workflow. AI-first regtech is the team I want to be operating in.

Attaching my resume and a 6-slide pitch deck mapping these proofs to the JD line-by-line, plus three specific 30-day experiments I'd run on Flagright's paid program. Full portfolio is at adhi-growth-ai.vercel.app/ref/flagright-pmm.

US remote. Looking forward to the conversation.

Portfolio: adhi-growth-ai.vercel.app/ref/flagright-pmm

Adhithya Rokhith Bhaskar