Adhithya Rokhith Bhaskar

San Francisco Bay Area, CA | arokhith@gmail.com | +1-929-563-9974

October 26, 2026

Hiring Team, Glean

When I built SwipeHire on the OpenAI API, the subscription price only worked once I knew what each user actually cost in model calls. I modeled three pricing tiers against that per-user API cost, then tested the result against willingness-to-pay signals from 12+ interviews before anyone paid for it. That is the same arithmetic behind a model hub: usage, cost and capacity have to pencil out before a customer ever gets to choose a model.

The rest of my product work has been on hardware, not language models, but the roadmap discipline carries over directly. I owned the roadmap and a 100+ feature backlog at Sensata, building the forecasting models that told Engineering which sensing capabilities to ship next, then running 30+ interviews with OEMs and distributors to check that plan against what they actually needed. The adoption curve for a new sensor capability and the upgrade path to a newer LLM are the same shape: someone has to make the case for switching, then measure whether the switch paid off.

At Applied Materials I build Claude agents to automate the repetitive side of market forecasting, and I review what they produce before it reaches a P&L number. Evaluating a model's output against a quality bar, provider by provider, is the same habit applied at a different scale, and it is the part of this role I would enjoy most.

Proof trail at adhi-product-ai.vercel.app/ref/glean.

Adhithya Rokhith Bhaskar