

Adhithya Rokhith Bhaskar

San Francisco Bay Area, CA | arokhith@gmail.com | +1-929-563-9974

October 2, 2026

Hiring Team, Glean

SwipeHire's auto-tailored resumes run on the OpenAI API, and the real work was never the API call. It was deciding what to do when the model's first answer was wrong. I built that review loop into the product from the PRD stage, zero-to-one, before I ever saw a roadmap that called it "AI Quality."

Most of my product experience lives on sensing hardware, not language models, but the job underneath is the same one your req describes. I owned the product roadmap and a 100+ feature backlog at Sensata, and that backlog only held up because I ran 30+ interviews with OEM and distributor engineers to find where the product actually broke, not where I assumed it would. Those fixes lifted CSAT 22% and feature adoption 18% across 50,000+ connected devices. That is customer-focused prioritization applied to a different kind of platform.

I also build my own Claude AI agents at Applied Materials to automate the repetitive side of forecasting, and I don't ship what they produce without reading it first. Build the automation, then build the check on the automation: that is the instinct I'd bring to a team deciding how good is good enough for an AI answer, especially one that sits in front of enterprise customers.

I like that this role gets measured by whether the answer was actually right, not just whether it shipped on time. Proof trail at adhi-product-ai.vercel.app/ref/glean.

Adhithya Rokhith Bhaskar