

PITCH DECK | SENIOR GROWTH MARKETER, HANDSHAKE AI

I built a smaller version of what you're building.

Now I want to scale it with you.

SwipeHire - AI career platform shipped solo, 50+ users, +25% activation

Pixis - \$1.4M aggregate book across 6 paid channels

NSE Bot - second AI product, Python + Claude API

4 Anthropic Claude Certifications

Cursor + Claude as daily workflow

Adhithya Rokhith Bhaskar

Boston, MA | arokhith@gmail.com | (929) 563-9974 | adhi-growth-ai.vercel.app/ref/handshake

SwipeHire and Handshake AI are the same product, at different scales.

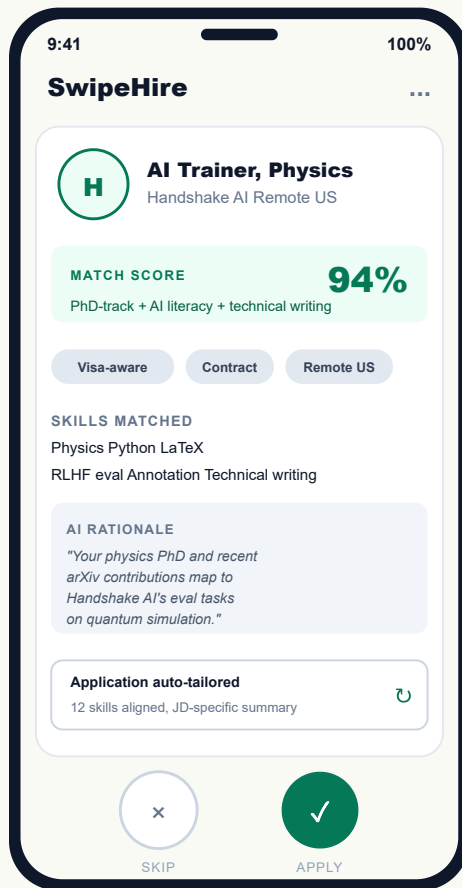
Most growth-marketing applicants pitch capabilities they could bring. I'd rather point at the product I've been running for the last two years that overlaps structurally with what Handshake AI is doing at \$1B scale.

DIMENSION	SwipeHire (mine, solo)	Handshake AI (yours, \$1B run-rate)
NETWORK ICP	Early-career + international students with specialized skill profiles - tomorrow's PhDs and researchers.	PhDs + researchers + subject-matter specialists contributing to frontier AI work.
MATCH LOGIC	LLM-scored JD-to-resume fit on top of deterministic skill overlap. Visa-aware filtering as a first-class signal.	Expert-to-task matching across evals, data labeling, model-eval workflows.
STACK	Python + OpenAI API + React. GA4 + GTM + Amplitude wired day one.	(Presumably) LLM orchestration over expert workflows, customer data warehouse, eval-loop infrastructure.
GROWTH MOTION	SEO + referral + paid social, A/B'd at low spend before scaling. Lifecycle email on activation + retention.	Multi-channel paid + SEO + community + lifecycle - the literal JD.
VALIDATION	50 users, 8 A/B tests on activation, +25% activation lift V1 to V2, weekly cohort retention.	30K experts, \$60M/month paid out, frontier-lab partnerships.

The operator surface is the same. What changes between SwipeHire and Handshake AI is scale, audience seniority, and the type of work the expert does at the other end. The eight-layer GTM motion - ICP, acquisition, onboarding, match, instrumentation, A/B, retention, community - is identical.

An AI career platform shipped as a GTM motion.

Python + OpenAI API + React. ICP defined, acquisition designed, onboarding rebuilt, A/B-tested, instrumented, retained. The eight layers, solo.



The eight layers, shipped solo

- 01 ICP defined.** Early-career + international students with specialized profiles. 12 user interviews before any code.
- 02 Acquisition.** SEO + referral + paid social, channel-tested at low spend before scaling.
- 03 Onboarding.** Resume parser on the OpenAI API, structured JSON output, skill normalization on a controlled taxonomy.
- 04 Match.** LLM-scored JD-to-resume fit. Deterministic skill overlap as the floor; LLM reasons on top.
- 05 Apply.** One-tap apply with a CV auto-tailored to the JD - skill alignment + JD-specific summary, generated server-side.
- 06 Instrument.** GA4 + GTM + Amplitude wired day one. Custom events on every meaningful step in the funnel.
- 07 Test.** 8 A/B variants on registration + activation flows. Sample-size + significance gates up front.
- 08 Retain.** Weekly cohort retention reviewed. Channel-by-cohort split surfaced where attribution lied.

Outcomes

50+

Early users onboarded

8

A/B tests on activation

+25%

Activation lift V1 to V2

4 wks

Of cohort retention tracked

Stack

Python

OpenAI API

React

GA4

GTM

Amplitude

The card above is mocked for Handshake AI's audience. The architecture - LLM match scoring, visa-aware filtering, auto-tailored apply - is the production build. Swap "Sr Marketing Analyst" for "AI Trainer, Physics" and the loop is identical.

Pixis: \$1.4M aggregate book across six paid channels.

Three years ago, on an 8-client D2C portfolio with per-client monthly budgets of \$15K-\$90K. The literal must-have in the JD, with the channel mix annotated.

\$1.4M AGGREGATE BOOK

illustrative split, 8-client D2C portfolio



PER-CLIENT MONTHLY BUDGETS, \$15K-\$90K 8 clients, weekly client CFO QBRs



Six channels. Eight clients. Weekly QBRs to client CFOs.
Holdout-incrementality work on one beauty client drove 3x ROAS in a single month.

Channel ownership (illustrative)

Google Ads	Branded + non-branded + Shopping	~30%
Meta	Prospecting + retargeting, creative cadence weekly	~26%
LinkedIn	ABM for B2B SaaS clients in the book	~16%
Criteo	Retail-media for D2C beauty + fashion clients	~15%
Amazon Ads	Sponsored Products + Brands for e-commerce	~11%
Programmatic	Adobe-driven DSP for awareness + retargeting	~2%

The diagnostic that mattered: blended ROAS read fine across the book. On the beauty client, a holdout-incrementality slice exposed two over-credited channels. Reallocated based on true lift; the brand cleared 3x ROAS the next month. That's the muscle Handshake AI's growth seat will need every week.

Three things I'd ship in week one for the expert network.

Concrete, not aspirational. Scoped to what would compound fastest for an expert-acquisition + retention motion at Handshake AI's current scale.

SHIP 1 - ACQUISITION

Top-50 CS-PhD ABM campaign on LinkedIn

LinkedIn ABM targeting active PhD candidates at top-50 CS departments. Audience layers built on program + research area + advisor. Creative tied to a specific eval-task domain. Lead form connected to a referral incentive paid on first contribution.

Why first. Universities are the deepest source of net-new researcher supply. ABM beats broadcast at PhD-level intent. The referral hook is the kind of structured experiment that pays for itself fast.

SHIP 2 - LIFECYCLE

"First Contribution in 14 Days" email sequence

Five-email lifecycle for newly-onboarded experts. Day 0 welcome with first task suggestion. Day 3, 7, 10, 14 nudges with progressively richer task matches. A/B'd on subject line, on send-time, on task-match algorithm. Drop-off measured at every step.

Why first. Activation is where expert-network economics live or die. A 14-day measured sequence cuts the time-to-first-contribution and gives the team a clean experimentation surface.

SHIP 3 - SEO + CONTENT

"Side income for researchers" content cluster

Five-piece content cluster optimized for "PhD AI training income", "academic side work in AI", "research RLHF side gigs", and adjacent intent queries. Schema marked up for AI Overview citation. Founder-led distribution via LinkedIn + Twitter + research forums.

Why first. Organic compounding. The expert network's recruiting story is a content category nobody owns yet. Owning these queries today pays back for years.

What I'd want from **the first fifteen minutes.**

What I'd want to understand: how Handshake AI is currently splitting growth investment between acquisition (university partnerships, paid ABM) and retention (lifecycle, expert satisfaction); which of the three week-one ships would die first inside the current stack and why; and where the team is putting their thumbs on the scale between SEO compounding and paid scaling right now.

What I'd bring: a one-pager on how I'd structure the first 90 days, an honest opinion on the tradeoff between optimizing for expert-network breadth (more contributors) versus depth (more contributions per contributor), and a working version of one of the week-one ships scoped to whatever sample data is shareable.

Remote works. Boston-based for whenever in-person matters.

15 min with the Handshake AI team

or written async if that's easier on your week

Adhithya Rokhith Bhaskar

arokhith@gmail.com | (929) 563-9974 | Boston, MA

Portfolio: adhi-growth-ai.vercel.app/ref/handshake

LinkedIn: linkedin.com/in/adhithyarokhith