

PITCH DECK | AI OPERATIONS, GTM

Building AI Operations for GTM.

Three proofs, not a pitch.

SwipeHire - AI product shipped solo, 50+ users, 8 A/B tests, +25% activation

NSE Bot - production analytics on the Claude API

Plug Power - MarTech from zero on HubSpot + Salesforce + Adobe

4 Anthropic Claude Certifications

Cursor + Claude as daily workflow

Adhithya Rokhith Bhaskar

Boston, MA | arokhith@gmail.com | (929) 563-9974 | adhi-growth-ai.vercel.app/ref/hightouch

Three things I've shipped **that map directly to this seat.**

Most application decks pitch capabilities. This one points at artifacts. Each of the three has a working version, an outcome someone read, and a direct line to the JD.

LEAD EXAMPLE

An AI product shipped as a GTM motion

SwipeHire

Job-matching product on Python + OpenAI API + React. Auto-tailors resumes to JDs and surfaces visa-aware matches. ICP defined, acquisition + onboarding designed, GA4 + GTM + Amplitude instrumented day one, 8 A/B tests run on activation, cohort retention tracked weekly. 50+ early users onboarded; activation up 25% from V1.

Why this matters here. *ICP, acquisition, build, instrumentation, A/B testing, retention. Every layer the AI Operations seat touches at Hightouch - just with job-seekers in the seat where SDRs and AEs will sit.*

A second agentic AI build

NSE Paper-Trading Bot

Python + Claude API end-to-end. Rule-based scoring (EMA + RSI + ATR) for the deterministic layer; Claude API for the agentic reasoning layer over news and catalyst context. Capital-constraint filter on top. Different vertical, same architecture pattern.

Why this matters here. *Two shipped AI products on two different LLM APIs (OpenAI + Claude) is rare at this seniority. It's the working proof, not the credential.*

GTM stack from zero


Plug Power MarTech

HubSpot + Salesforce + Adobe + GA4 + GTM, architected end-to-end. UTM spec written from zero. Sequences, pipelines, lifecycle stages, account-level segmentation. Closed-loop attribution from ad click to Sales-Qualified Lead. CPA fell 25%; conversion rose 20%.

Why this matters here. *"Audit GTM workflows and identify automation opportunities" is the literal description of this work.*

An AI product shipped as a GTM motion.

SwipeHire matches job-seekers to roles via an LLM-scored swipe-card UI, auto-tailors a CV to the JD on apply, and is visa-aware out of the box. Built solo on Python + OpenAI API + React. Below: a working screen, the eight layers shipped, and the instrumentation that made the A/B tests honest.



The eight layers, shipped solo

- 01 ICP defined.** Early-career + international students, visa-aware. Sized via 12 user interviews before any code.
- 02 Acquisition.** SEO + referral + paid social, channel-tested at low spend before scaling.
- 03 Onboarding.** Resume parser on the OpenAI API, structured-output JSON, skill normalization on a controlled taxonomy.
- 04 Match.** LLM-scored JD-to-resume fit. Deterministic skill overlap as the floor; LLM reasons on top, not in place of, the score.
- 05 Apply.** One-tap apply with a CV auto-tailored to the JD - 8 keyword alignments, JD-specific summary, generated server-side.
- 06 Instrument.** GA4 + GTM + Amplitude wired day one. Custom events on every meaningful step in the funnel.
- 07 Test.** 8 A/B variants run on registration + activation flows. Sample-size + significance gates set up front.
- 08 Retain.** Weekly cohort retention reviewed. Channel-by-cohort split surfaced where attribution lied.

Outcomes

50+

Early users onboarded

8

A/B tests on activation

+25%

Activation lift vs V1

4 wks

Of cohort retention tracked

Stack

Python

OpenAI API

React

GA4

GTM


Amplitude

Hightouch translation. This is the same eight-layer loop AI Operations runs internally - just with SDRs/AEs/CS in the seat where job-seekers sit. The instrumentation discipline transfers directly; the LLM-scored matching becomes lead/account scoring; the swipe-card UI becomes a Slack-delivered "next best action" card.

Hightouch | Adhithya Rokhith Bhaskar

The agentic layer pattern, on a different vertical.

NSE Paper-Trading Bot architecture. The same "deterministic floor, LLM ceiling" pattern that powers SwipeHire's matching, expressed on financial-market data instead of jobs.



TWO BUILDS, ONE PATTERN

SwipeHire: deterministic skill-overlap score + OpenAI reasoning + UI for human-in-the-loop validation (swipe).

NSE Bot: deterministic technical score + Claude reasoning + dashboard for backtest validation.

Same shape, different vertical, different LLM provider. The pattern is what transfers to Hightouch's agentic-workflow seat.

*Deterministic floor keeps the agentic layer honest. Claude reasons on top of structured signal, not in place of it.
That separation is what makes the output trustworthy at production scale - same principle as SwipeHire's match score.*

What GTM stack work looks like **in my hands.**

Plug Power MarTech architecture. Built from zero before "GTM automation" became the buzzword on every JD.

CHANNELS

Paid Acquisition

Google Ads, LinkedIn, Meta, Adobe-driven programmatic - A/B on keywords, audiences, creative

TRACKING

Instrumentation Layer

UTM spec written from zero | GTM containers rebuilt | GA4 events validated end-to-end | Adobe Analytics

AUTOMATION

HubSpot - architected, not inherited

Sequences | Pipelines | Lifecycle stages | Account-level behavioral segmentation | MQL nurture drip

HANDOFF

Salesforce

Campaign attribution validated end-to-end | Lead routing rules | SLA documentation | Closed-loop to SQL

OUTCOMES THAT MATTERED

-25% CPA, +20% CVR are the numbers most readers stop at. The work I would point to is the layer underneath.

UTM spec from zero. GTM containers rebuilt. Lead routing validated through Salesforce campaign attribution.

Different stack than Clay or Hightouch CDP. Same operator seat.

-25%

CPA across channels

+20%

Conversion rate

From 0

MarTech architected

4-vendor

RFP run end-to-end

Three things I'd ship in week one.

Concrete, not aspirational. Scoped to where I'd start day one given what's public about Hightouch's stack and what the JD lists as nice-to-have.

SHIP 1 - CLAUDE SKILL

"AE Account Brief" skill

Takes a Salesforce account ID. Pulls open opps, recent CS activity tagged via Hightouch CDP, last earnings call summary or 10-K diff via Glean. Output: a one-pager that lands in Slack 30 minutes before the call.

Why first. AEs already do this manually. A Claude skill cuts the prep from 20 minutes to 3 with better consistency. Fast first ship; clean enablement story to the GTM org.

SHIP 2 - WORKFLOW AUDIT

Clay enrichment hit-rate vs Hightouch CDP

Compare current Clay ICP enrichment hit-rate against composable CDP segments. Find the 20% of accounts where switching enrichment source compounds the loop. Document where each stack wins.

Why first. Dogfooding. RevOps teams trust a story they can tell with their own product more than they trust slideware. This is also the kind of audit the JD's first responsibility line describes.

SHIP 3 - KNOWLEDGE RESOURCE

"Why We Won" library in Glean

Tag the last 100 closed-won opportunities by competitor displaced, objection-handling pattern that closed, and the deal that was replaced. Load to Glean as a knowledge resource. Wrap in a Claude skill for SDR objection prep.

Why first. Compounding artifact - every new deal feeds it. It's the kind of thing GTM teams nod about but never actually build, because no one owns it. AI Operations does.

What I'd want from **the first fifteen minutes.**

What I'd want to understand: how AI Operations splits today between RevOps, Solutions Engineering, and the AI/agents team; which of the three week-one ideas would die first inside Hightouch's stack and why; and where the team would put their thumbs on the scale for ICP enrichment dogfooding versus customer-facing builds.

What I'd bring: a one-pager on how I'd structure the first 90 days, an honest opinion on the prototype-vs-platform tradeoff for early agentic deployments, and a working version of the AE Account Brief skill scoped to whatever sample data is shareable. Plus a SwipeHire walk-through, if useful.

Remote works. Boston for whenever in-person matters.

15 min with the RevOps + AI team

or written async if that's easier on your week

Adhithya Rokhith Bhaskar

arokhith@gmail.com | (929) 563-9974 | Boston, MA

Portfolio: adhi-growth-ai.vercel.app/ref/hightouch

LinkedIn: linkedin.com/in/adhithyarokhith