Adhithya Rokhith Bhaskar

San Francisco Bay Area, CA | arokhith@gmail.com | +1-929-563-9974

October 25, 2026

Hiring Team, Rubrik

Your posting talks about moving the CPQ rules engine toward AI agents with human-in-the-loop checkpoints. I already built that machine. PriceKeel's guardrail engine recommends an explainable discount guardrail learned from a company's own closed-won and closed-lost deals, and every human override gets logged in an audit trail Finance can check on demand. That is not a line item about agentic AI, it is the product I run, including the CRM, CPQ and deal-approval workflow underneath it.

The quote-to-cash half of the role is where I spent the year before PriceKeel. At Sensata I repriced parts sold through distribution and raised distributor and channel-partner margins from 53% to 61% in under two quarters, negotiated multimillion-dollar customer quotes directly with senior executives, and built the FY2026 Bleeders and Leakers Power BI program: 1,200 pricing opportunities tracked against a $627M target across 3 regions and 20 product lines, surfacing $433M in margin gaps for leadership to prioritize. That fed a $700K+ margin recovery in two quarters. I have simplified multi-product, multi-channel quoting from the inside, not in theory.

Before pricing, I owned the actual product roadmap: a 100+ feature backlog for sensor solutions, run with Engineering in Agile across 3 regions, cutting time-to-market 20%. I am Strategic Marketing III at Applied Materials now, automating repetitive pricing and quoting analysis with AI agents built on Claude, so the agentic half of your roadmap is something I already run day to day.

Rubrik's quoting operating model has to serve Product, Sales, Finance, Deal Desk and channel leaders at once while the rules engine gets rebuilt into something agentic, without losing anyone's trust in the exception trail. That is the exact problem PriceKeel's guardrail engine was built to solve. Proof trail at adhi-pricing-ai.vercel.app/ref/rubrik.

Adhithya Rokhith Bhaskar