

MARKETING ANALYST | SAMPLE WORK

Mapping the JD to work I have shipped.

For Speak. Designed to show, not describe.

Each slide answers one of your four JD bullets with concrete past work.

Adhithya Rokhith Bhaskar

Boston, MA | arokhith@gmail.com | (929) 563-9974

Portfolio: adhi-growth-ai.vercel.app/ref/speak | LinkedIn: [/in/adhithyarokhith](https://in/adhithyarokhith)

Dashboards leadership reads on Mondays.

"Develop analyses, dashboards and reporting to measure the effectiveness of different marketing strategies."

PROGRAM TRACKED

\$627M

margin-recovery opportunity

RECORDS

1,200

pricing opportunities

COVERAGE

3 / 4 / 20

regions / families / lines

ADOPTION LIFT

+18%

on 50K+ devices (Amplitude + GA4)

SENSATA MARGIN-RECOVERY DASHBOARD | BLEEDERS VS LEAKERS BY REGION

REGION	BLEEDERS (\$M)	LEAKERS (\$M)	RECOVERED	STATUS
NA	142.0	68.5	+62%	on track
EU	98.4	51.2	+47%	on track
APAC	76.1	39.0	+18%	at risk
LATAM	58.8	22.5	+41%	on track

WIN-RATE TREND BY PRODUCT FAMILY (LAST 4 QUARTERS)



WHAT I BUILT AND OWN

- **Sensata \$627M Power BI:** 4-page executive view tracking 1,200 opportunities. Senior leadership reads it weekly.
- **50K+ device KPI framework:** SQL + GA4 + **Amplitude** for product-feature engagement tracking. Adoption +18%.
- **Pixis 8-client portfolio reporting:** per-client ROI and iROAS dashboards on \$1.4M media spend. Weekly QBR cadence.
- **Plug Power closed-loop view:** click-to-SQL one number rather than three. Daily attribution across paid + email + organic.

The dashboard is the strategy artifact. The weekly review is the operating rhythm. Both have to exist for either to work.

Activation event > activation rate.

"Develop and improve internal models for things like attribution, media mix, and LTV forecasting."

ACTIVATION EVENT

3+ AI sessions / W1

predicts D30 retention

D30 RETENTION | ACTIVATED

76%

vs 6% non-activated, 31% blended

ACTIVATION LIFT SHIPPED

+25%


via 8 A/B tests on the funnel

RENEWAL MODEL ACCURACY

87%

on 30-day signal at Pixis

SWIPEHIRE D30 RETENTION CURVE | THREE COHORTS



Method: Identified the activation event (three or more AI tutor conversations completed in week one) by regressing D30 retention on event-stream variables in Amplitude. Instrumented the onboarding funnel to optimize for the event, then validated lift via eight A/B tests on activation flows. The retention gap (76% vs 6% on D30) is what makes this LTV-load-bearing rather than vanity.

PIXIS B2B SAAS RENEWAL PROBABILITY | BY D30 SIGNAL

Engagement Segment (D30)	N	Renewal
QBR engaged + 3+ campaigns + ROAS > 2.5x	2	94%
QBR engaged + 1 to 2 campaigns + ROAS > 2x	3	78%
QBR engaged + minimal campaign activity	1	61%
QBR ghosted + 1 to 2 campaigns	1	42%
QBR ghosted + 0 campaigns + ROAS slipping	1	11%

INTERVENTION TRIGGERED WHEN PROBABILITY < 50%

Action	Lift
Creative refresh + channel re-key	+34%
Replace underperforming channel	+27%
Per-client renewal-aligned playbook	+19%

Churn -50%


\$1.4M book retained

8-person team

Closed loop made every model trustable.

"Develop and improve internal models for things like attribution, media mix, and LTV forecasting."

PLUG POWER B2B PAID MIX | AFTER ATTRIBUTION REBUILD



- Paid Search** | 32% | Google Ads, SA360
- Paid Social** | 24% | Meta, LinkedIn
- Programmatic** | 18% | Adobe Audience Manager
- Email + lifecycle** | 14% | Adobe + Salesforce
- Organic** | 12% | SEO + content

CPA

-25%

post-rebuild

CONVERSION

+20%

click-to-SQL

STACK STITCHED INTO ONE NUMBER

- Wrote **UTM spec from scratch**
- Rebuilt **GTM containers**
- Validated lead routing into **Salesforce campaign attribution**
- **Adobe Audience Manager** for account-level segmentation + suppression
- **GA4** for funnel + conversion event tracking
- Reconciled platform numbers vs SQL data on weekly attribution review

MMM-ADJACENT ECONOMETRICS


Built an **R-based regression model** for demand elasticity at Sensata. Variable selection, hypothesis testing, cross-validation, interaction-term interpretation. Forecasts profit-maximizing price points across feature scenarios. Same econometric muscle MMM uses, applied to pricing.

Most "attribution" failures are tagging failures. Fix the plumbing before you fix the model. Once the data was clean, every model on top of it became trustable.

Holdout proved the lift was causal.

"Design and orchestrate incrementality tests to identify stronger causal relationships between campaigns and conversion."

PIXIS FASHION D2C PARTNER | GEO + CREATIVE-VARIANT HOLDOUT (30 DAYS)



Test design: Geo + creative-variant holdout. Treatment regions received Shopping feed cleanup, smart bidding, creative refresh, lookalike expansion. Control regions ran baseline. Reconciled the lift against creative-variant diagnostics that aggregate ROAS would have hidden.

TEST DESIGN CHECKLIST

ELEMENT	DECISION
Hypothesis	Lift > 1.5x is causal, not seasonal
Unit of randomization	Geo (DMA-level)
Sample size	Pre-registered, power 80%
Treatment	Full-funnel rebuild
Control	Matched geo, baseline mix
Duration	30 days, weekly check-in
Read-out	Lift + creative-variant diagnostics

OTHER TESTS SHIPPED

- **Plug Power:** lift-testing methodology for net-new MQL evaluation, separating causal lift from correlation in a B2B funnel.
- **SwipeHire:** 8 structured A/B tests on activation flows. Pre-registered metrics + sample size. Win/loss/no-conclusion taxonomy documented.
- **Pixis:** written measurement playbooks (metric defs + lift cadence + reporting standards).

Stack you build with is stack I build with.

Marketing analytics is a tool problem when the tools fight you, and a craft problem when they don't. Speak's stack is one I already work in.

PRODUCT ANALYTICS

Amplitude (event taxonomy, retention cohorts, funnel, A/B), GA4, Google Tag Manager, Adobe Analytics

ATTRIBUTION + MODELING

Salesforce, Adobe Audience Manager, UTM governance, MTA, R regression (MMM-adjacent), cohort & survival

REPORTING

Power BI, Tableau, SQL (advanced), Python (pandas, numpy), Excel

AI-FLUENT EXECUTION

OpenAI API, Claude API, 4 Anthropic Claude certs, prompt-based iteration as actual workflow

Why Speak, specifically

SwipeHire shipped on OpenAI API. Four Claude certs come from running prompt-based iteration as my actual workflow, not as a side project. AI-fluent measurement is not a learning curve. Speak is OpenAI-backed. Cultural-fit signal is unusually clean.

Why this role, specifically

The four JD bullets are the four problems I have spent the last four years solving. Alexandra came from AppLovin where the bar for marketing analytics is set by holdout-validated UA, cohort LTV, and AppDiscovery-grade attribution rigor. That is the bar I want to work to.