

PITCH DECK | GROWTH MARKETING MANAGER, AI

AI as an operational foundation.

The shape I've already shipped twice. Solo.

SwipeHire - AI product on Python + OpenAI API + React

NSE Bot - second AI product on Claude API

4 Anthropic Claude Certifications

8 enterprise marketers trained on AI tools at Pixis

Cursor + Claude as daily workflow

Adhithya Rokhith Bhaskar

Boston, MA | Remote | arokhith@gmail.com | (929) 563-9974 | adhi-growth-ai.vercel.app/ref/suite-studios


The "AI as operational foundation" line is the rarest one in your JD.

Most growth marketers have used AI tools. Few have shipped AI products solo, end-to-end, with a feedback loop they own. The JD asks for that exact shape. Here's where mine maps.

JD MANDATE	Suite Studios asks for	What I've shipped
AI PRODUCTS	Hands-on AI tools for research, content, workflow automation, performance analysis.	Two shipped AI products solo. SwipeHire (Python + OpenAI API + React; LLM-scored matching; GA4 + GTM + Amplitude day-one). NSE Bot (Python + Claude API as agentic reasoning over rule-based scoring).
REPEATABLE WORKFLOWS	Develop repeatable AI-driven workflows for marketing operations.	SwipeHire's eight-layer GTM loop (ICP -> acquisition -> onboarding -> match -> apply -> instrument -> test -> retain) is the working model. Same shape as any marketing-ops AI workflow at production scale.
OPERATIONAL FOUNDATION	Establish AI as an operational foundation across the marketing function.	Trained 8 enterprise marketers on AI bidding + creative tools at Pixis through weekly client CFO QBRs. Structured curriculum, sequenced what to teach, measured what stuck.
CONTENT + LAUNCHES	Content, email, website, launches, customer narratives. Translate product capabilities into clear messaging.	Sensata customer-story pipeline (30+ interviews -> positioning + 50-rep sales enablement library -> +30% win rate). Plug Power AEO/GEO program (FAQ pages structured for ChatGPT/Claude/Perplexity citation).
WEBSITE + LIFECYCLE	Manage website content, landing pages, conversion optimization.	Plug Power HubSpot lifecycle architected from

The eight-layer AI workflow, shipped solo.

SwipeHire is the working version of what the JD's "repeatable AI-driven workflow for marketing operations" line describes - just with job-seekers as the audience instead of brand marketers.



The eight-layer loop, shipped solo

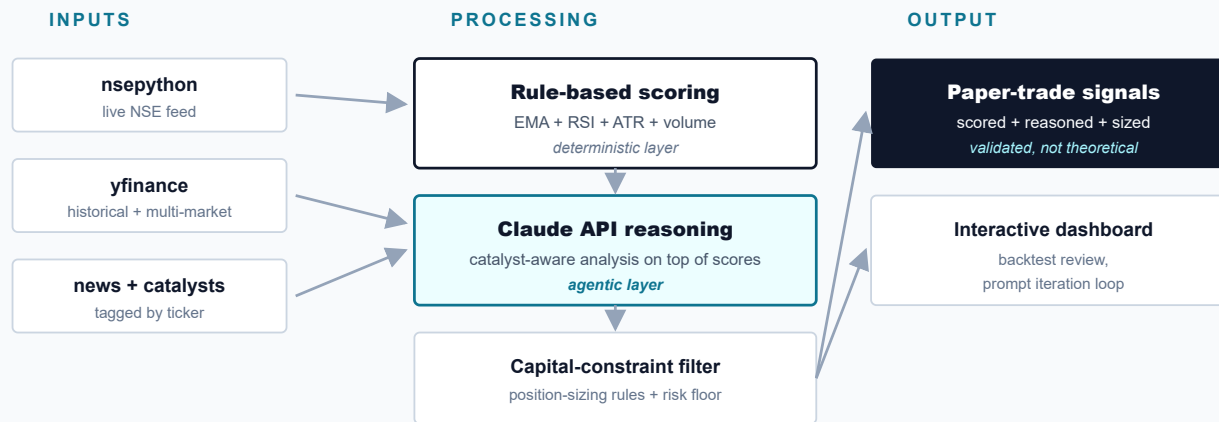
- 01 ICP defined** from 12 user interviews before any code.
- 02 Acquisition** channels A/B'd at low spend before scaling.
- 03 Onboarding.** Resume parser on the OpenAI API, structured JSON output, skill taxonomy normalization.
- 04 LLM matching.** Score on deterministic skill overlap first; LLM reasons on top.
- 05 Auto-tailored apply.** JD-specific CV generation server-side.
- 06 Instrument.** GA4 + GTM + Amplitude wired day one.
- 07 Test.** 8 A/B variants on activation; +25% lift V1 -> V2.
- 08 Retain.** Weekly cohort retention reviewed line by line.

Suite Studios translation

Replace "job-seekers" with "creative pros at production studios." Replace "JD" with "Suite product capability." Replace "match score" with "feature-fit recommendation." The loop is the same.

The agentic-layer pattern, a second time.

NSE Paper-Trading Bot architecture. Same "deterministic floor, agentic ceiling" shape as SwipeHire - just on financial data instead of job data. Same shape any marketing-ops AI workflow takes at production scale.



Deterministic layer keeps the agentic layer honest. Same pattern would govern any marketing-ops AI workflow (content gen, lifecycle copy, AEO audit).

TWO BUILDS, ONE PATTERN

SwipeHire: deterministic skill overlap + OpenAI reasoning + measured user feedback (swipe). **NSE Bot:** deterministic technical score + Claude reasoning + measured outcome (backtest).

Same shape. Different vertical. Different LLM provider. That pattern is what transfers to "AI as operational foundation" at Suite.

Three AI workflows I'd ship in the first thirty days.

Concrete and scoped, not aspirational. Each one is a repeatable AI-driven marketing-ops workflow with measurable output - exactly what the JD asks for.

WORKFLOW 1

Customer-Story-to-Content pipeline

Claude pipeline that takes a customer interview transcript and outputs drafts: 1-page case study, three LinkedIn posts, a launch FAQ page (schema-marked for AI citation), and three email-sequence variants. Reviewer-in-the-loop, not auto-publish.

Why first. The JD names "customer storytelling" + "translate product capabilities" + "AI-driven workflows for marketing ops" in three bullets. This workflow hits all three. Sensata 30-interview pipeline is the analog I built manually; this is the AI-accelerated version.

WORKFLOW 2

AEO Audit on Launch

Pre-launch + post-launch: Claude-driven comparison of Suite's brand visibility on ChatGPT, Claude, Perplexity, and Google AI Overviews for ~20 high-intent queries (creative collab tools, post-production workflow, AI for editorial). Output: weekly delta + content-gap action list.

Why first. Launches without AEO measurement leave half the visibility on the table. Plug Power's AEO/GEO program is the proof I've built this loop before, manually. Suite is the right place to productize it.

WORKFLOW 3

Lifecycle Email Personalization at Variant Scale

Cursor + Claude pipeline that takes a single source brief and generates segment-specific email-copy variants (editor, producer, studio exec) plus a draft A/B test plan with sample-size + significance gates. Output: ready-to-load HubSpot sequences.

Why first. Lifecycle email is where most B2B marketing teams write the same copy six times manually. This workflow compresses six versions into one prompt session, with the deterministic layer (segment definitions, send rules) keeping it honest.

What I'd want from the first fifteen minutes.

What I'd want to understand: where the marketing function currently sits between "uses AI tools" and "operationalized AI workflows" - and which gap matters most. Which of the three week-one workflows would die first inside Suite's stack and why. And where the team would put their thumbs on the scale for content compounding versus launch-velocity right now.

What I'd bring: a one-pager on how I'd structure the first 90 days, a draft Customer-Story-to-Content pipeline scoped to one of Suite's existing customer narratives, and an honest take on which AI tools (Claude vs GPT vs OSS) I'd use where and why.

Remote works. Boston-based for whenever in-person matters.

15 min with the Suite Studios team

or written async if that's easier on your week

Adhithya Rokhith Bhaskar

arokhith@gmail.com | (929) 563-9974 | Boston, MA

Portfolio: adhi-growth-ai.vercel.app/ref/suite-studios

LinkedIn: linkedin.com/in/adhithyarakhith