

Adhithya Rokhith Bhaskar

San Francisco Bay Area, CA | arokhith@gmail.com | +1-929-563-9974

September 22, 2026

Hiring Team, Turo

At PriceKeel, the company I founded, I built a guardrail engine that reviews a live B2B pricing exception and returns an explainable recommendation learned from a company's own closed-won and closed-lost deals, then keeps an audit trail so the person looking at it can see exactly why the system said what it said. That is the same problem your Host Pricing team is solving at a different scale: a host looking at a suggested price needs to understand it before they will act on it, and the moment they cannot, they override it and both sides lose a little.

I have built that kind of trust into a pricing tool before, not just a model that scores well on paper. At Sensata I ran the Bleeders and Leakers program, a pricing recommendation system that scored 1,200 opportunities against a \$627M target and told Sales and Finance not just what to change but why, and it recovered \$700K+ in margin in two quarters because people actually acted on what it recommended instead of shelving another dashboard. Getting a recommendation adopted is a different problem than getting it accurate, and it is the one I keep ending up in charge of.

Before pricing became the focus, I owned the product roadmap and a 100+ feature backlog for Sensata's sensor products, running Agile with a cross-functional engineering team across three regions with no existing playbook to inherit. Thirty-plus customer interviews told me which tools people trusted and which ones they quietly worked around, and closing that gap is what took CSAT up 22% and feature adoption up 18%. I would bring the same habit, watch what people actually do with a recommendation, not just what they say about it, to how hosts price their cars.

I would like to talk about what is actually driving host overrides today and where the tradeoffs are hardest to call. Proof trail at adhi-pricing-ai.vercel.app/ref/turo.

Adhithya Rokhith Bhaskar